

1

Data exploration

SUPPLY AND DEMAND, chocolate and peanut butter, education and income . . . some things just go hand in hand. Empirical relationships like these provide the foundation that allows us to build theories of how the world works.

Supply and demand: Economic theory and common sense predict that people will demand less of a thing as its price rises. How are we to explain, then, why demand for certain goods—a Rolls Royce, a Gucci handbag, a bottle of Dom Perignon—seems to increase with price, in apparent defiance of gravity? The empirical price–demand relationship for these *Veblen goods* requires a causal story, the canonical one being that people buy luxury items as status symbols.

Veblen goods are named after economist Thorstein Veblen. Another example is provided by so-called *Giffen goods*, where the usual price–demand relationship is reversed because of something called the income effect.

Chocolate and peanut butter: Sometimes empirical relationships can suggest new theories. For example, one notable culinary fact is that people, regardless of culture, tend to like sweet and salty things together: melon and ham; coconut and curry; and yes, Reese’s pieces. How does this fact influence the research agenda of scientists who study the physiology of human taste buds? Might sodium play a role in the perception of every basic taste, and not just the taste of salty things?

If you like coffee but don’t like the bitter taste, try adding a wee pinch of salt for every 2 tablespoons of coffee grounds, all mixed together when dry, next time you brew a pot. . . .

Education and income: “Human capital” sounds like a phrase of recent vintage. But it dates at least to Adam Smith:

The acquisition of such talents, by the maintenance of the acquirer during his education, study, or apprenticeship, always costs a real expense, which is a capital fixed and realized, as it were, in his person. Those talents, as they make a part of his fortune, so do they likewise that of the society to which he belongs. The improved dexterity of a workman may be considered in the same light as a machine or instrument of trade which facilitates and abridges labor, and which, though it costs a certain expense, repays that expense with a profit.¹

¹ *An Inquiry into the Nature And Causes of the Wealth of Nations*, Book 2.

The natural empirical question is: if education is a capital investment, how good is the return?²

In each case, a particular idea about how things work turns upon the interpretation of an observed relationship between things we can measure. To do this correctly requires care, judgment—and the right toolkit. The goal of this chapter is to equip you with some basic visual and numerical tools for exploring multivariate data sets, with an eye towards finding interesting relationships.

Cases and variables. In statistics, we typically refer to the *cases* and *variables* of a data set. The cases are the basic observational units that we’re interested in: people, houses, cars, guinea pigs, etc. The variables are the different kinds of information we have about each case—to cite an example we’ll soon come to, the horsepower, fuel economy, and vehicle class for several hundred cars.

Variables come in two basic kinds: numerical variables are represented by a number, while categorical variables are described by the answer to a multiple-choice question. This chapter will describe some strategies for summarizing relationships among both kinds of variables, as well as some further refinements to this basic “numerical versus categorical” distinction.

Variation across categories

MUCH OF the data you’ll meet will involve categories: man or woman; chocolate or vanilla; Toyota, Honda, or Hyundai; butcher or baker or candlestick maker. A simple, effective way to summarize these categorical variables³ is to use a *contingency table*. On the Titanic, for example, a simple two-way table reveals that women and children survived in far greater numbers than adult men:

	Girl	Woman	Boy	Man
Survived	50	242	31	104
Died	22	74	51	472

We call this a two-way or bivariate table because there are two variables are being compared: survival status versus type of person. The categories go along the rows and columns of the table; the cell counts show how many cases fall into each class. The pro-

² This is a fascinating topic for further study, to which a good starting point is [this article](#) from the St. Louis branch of the Federal Reserve.

³ Categorical variables are sometimes referred to as *factors*, and the categories themselves as the *levels* of the factor. The R statistical software package uses this terminology.

Table 1.1: A two-way table, because there are two categorical variables by which cases are classified. The data are available in the R package *effects*. Originally compiled by Thomas Cason from the *Encyclopedia Titanica*.

cess of sorting cases into the cells of such a table is often called *cross-tabulation*.

We can also make multi-way tables that show more than two variables at once. Given the constraints of a two-dimensional page, multiway tables are usually displayed as a series of two-way tables. As the following three-way table reveals, richer passengers, of either sex, fared better than others.

Cabin Class		1st	2nd	3rd
Female	Survived	139	94	106
	Died	5	12	110
Male	Survived	61	25	75
	Died	118	146	418

Table 1.2: An example of a *multi-way table*, where counts are classified by cabin class, sex, and survival. NB: passengers of unknown age are included in this table, but not the previous one.

Tables are almost always the best way to display small data sets with few classifying variables, for the simple reason that they convey a lot of information in a small space.⁴

Ordinal and binary variables. If a categorical variable has only two options (heads or tails, survived or died), we often call it an indicator, binary, or dummy variable. (These names can be used interchangeably.)

Some categories have a natural ordering, like measures of severity for a hurricane, or responses to a survey about consumer satisfaction. (Has your experience with our call center been Atrocious, Merely Bad, Acceptable, Good, or Excellent?) These are called *ordinal variables*. Ordinal variables differ from numerical variables in that, although they can be placed in a definite order, they cannot be compared using the laws of arithmetic. For example, we can't subtract "Good" from "Excellent" and get a meaningful answer, while we can subtract \$1000 from \$5000.

Mosaic plots. Categorical data can also be displayed in a *mosaic plot*. The area of each box tells you what fraction of cases fall into the corresponding cell of the contingency table. At a glance, this gives you an idea of how category membership predicts some of the variation in the outcome variable. Compare Figure 1.1 with the previous table of Titanic survivors stratified by sex.

Mosaic plots are at their best when used for data sets with more than two cross-classifying variables. In this context they are more

⁴ This animation provides some good guidelines for formatting tables.




Figure 1.1: Survival of passengers on the Titanic, stratified by the categorical predictor “girl/boy/man/woman.”

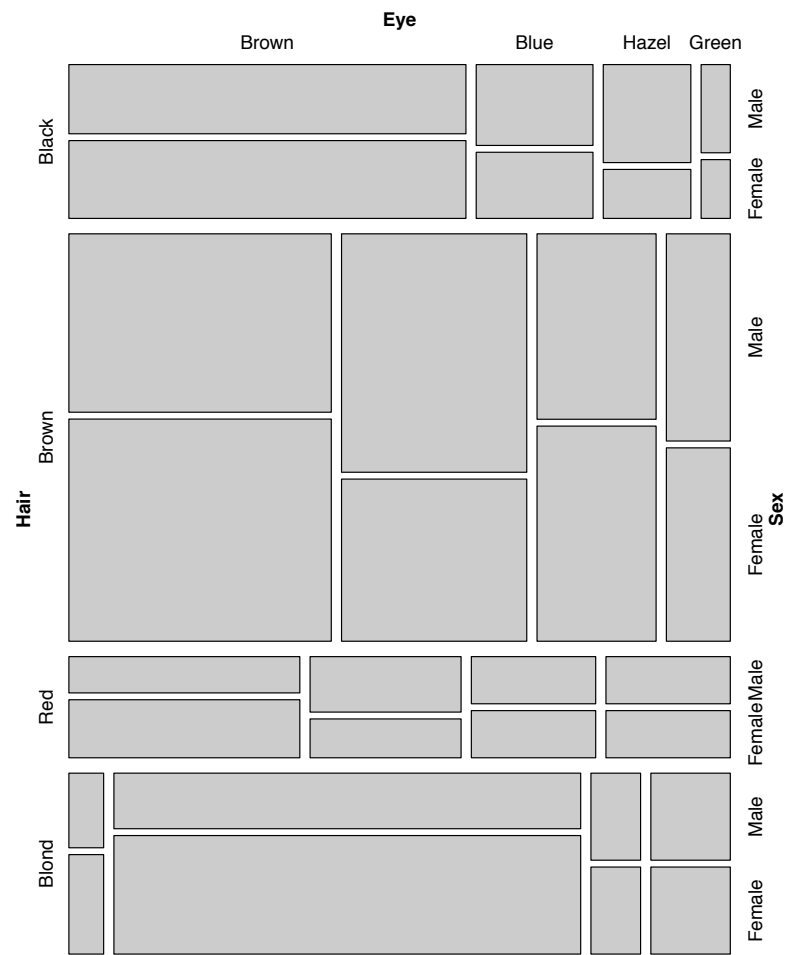


Figure 1.2: Hair color, eye color, and sex for 592 students at the University of Delaware. Data from Snee (1974), “Graphical display of two-way contingency tables.” *The American Statistician*, 28, 9–12. It appears as though, among people with blond hair and blue eyes, women are over-represented. (Compare the sizes of the male and female boxes for the blond-haired, blue-eyed crowd in the lower left.) It also appears as though green eyes are rarest among people with black hair, in a relative sense. (Notice how the “Green eyes” boxes are narrowest left to right in the row for black hair, compared to the row for any other hair color.)

useful for quick visual comparisons than are tables. For example, Figure 1.2 depicts joint variation in hair color, eye color, and sex for a sample of 592 students at the University of Delaware. Try answering the following questions using the plot.

- (1) Among students with blond hair and blue eyes, who is over-represented: men or women?
- (2) For which hair color would green eyes be the most surprising?

Relative risk and odds ratio

The **relative risk**, sometimes also called the risk ratio, is a widely used measure of association between two categorical variables. To introduce this concept, let's examine a tidbit of data from the PREDIMED trial, a famous study on heart health conducted by Spanish researchers that followed the lifestyle and diet habits of thousands of people over many years, beginning in 2003.⁵

The main purpose of the PREDIMED trial was to assess the effect of a Mediterranean-style diet on the likelihood of someone experiencing a major cardiovascular event (defined by the researchers as a heart attack, stroke, or death from cardiovascular causes). But as part of the study, the researchers also collected data on whether the trial participants were, or had ever been, regular smokers. The table below shows the relationship between smoking and whether someone experienced a cardiovascular event during the study period.

	Current or former smoker?	
	No ($n = 3892$)	Yes ($n = 2432$)
No event	3778	2294
Event	114	138

Let's compare the risk of cardiovascular events for smokers, versus that of non-smokers.⁶ Among the smokers, 138 of 2432 people (5.7%) experienced an event; while among the non-smokers, 114 of 3892 people (2.9%) experienced an event. To compute the relative risk of cardiovascular events among smokers, we take the ratio of these two quantities:

$$\text{Relative risk} = \frac{138/2432}{114/3892} = 1.94.$$

This ratio says that smokers were 1.94 times more likely than non-smokers to experience a cardiovascular event during the study.⁷

⁵ Estruch R, Ros E, Salas-Salvado J, et al. Primary prevention of cardiovascular disease with a Mediterranean diet. *N Engl J Med* 2013;368:1279-1290. The full text of the article is available at <http://www.nejm.org/doi/full/10.1056/NEJMoa1200303>

⁶ By "risk," we simply mean the chance of an event happening.

⁷ Of course, this doesn't prove that the smoking caused the cardiovascular events. One could argue that the smokers may have had other systematically unhealthier habits that did them in instead, and the smoking was merely a marker of these other habits. We'll soon talk about this issue of confounding much more.

More generally, for any event (a disease, a car accident, a mortgage default) and any notion of “exposure” to some factor (smoking, driving while texting, poor credit rating), the relative risk is

$$\text{Relative risk} = \frac{\text{Risk of event in exposed group}}{\text{Risk of event in non-exposed group}}.$$

The relative risk tells us how much more (or less) likely the event is in one group versus another. It’s important to remember that the relative risk (in our example, 1.94 for smokers) is quite different from the *absolute risk* (in our example, 0.057 for smokers). This distinction is often missed or elided in media coverage of health issues. See, for example, [this blog post](#) from the UK’s cancer-research funding body about news reports of cancer studies.

Odds ratio. An alternative measure of association in two-way tables is the *odds ratio*. This is defined as follows:

$$\text{Odds ratio} = \frac{\text{Odds of event in exposed group}}{\text{Odds of event in non-exposed group}}.$$

Recall that the odds of any event is itself a ratio:

$$\text{Odds of event} = \frac{\text{Chance of event happening}}{\text{Chance of event not happening}}.$$

This is what makes the odds ratio a bit confusing at first: it’s a ratio of ratios.

Let’s work out the odds ratio of cardiovascular events for smokers versus non-smokers from the PREDIMED study. For smokers, the odds of an event are

$$\text{Odds of event for smokers} = \frac{138/2432}{2294/2432} = \frac{138}{2294} \approx 0.06.$$

And for non-smokers:

$$\text{Odds of event for non-smokers} = \frac{114/3892}{3778/3892} = \frac{114}{3778} \approx 0.03.$$

Therefore the odds ratio of events for smokers is

$$\text{Odds ratio for smokers} = \frac{138/2294}{114/3778} \approx 1.99.$$

Most people find the odds ratio a bit tricky to interpret. Thus a natural question is: when should one use the odds ratio, as opposed to the relative risk? The best answer is that for certain types of complicated statistical models, measuring association using the odds ratio makes the math and computation much more convenient. We’ll see this issue return when we learn a technique called logistic regression, much later in the book.




Figure 1.3: Daily average temperatures for San Diego and Rapid City, 1995–2011. These data are visualized in a histogram, which is a simple and effective way to depict the variation of a single numerical variable across many cases. To construct a histogram, we first partition the range of possible outcomes (here, temperatures) into a set of disjoint intervals (“bins”). Next, we count the number of cases that fall into each bin. Finally, we draw a series of rectangles over each bin whose height is equal to the count within each bin. (Technically this is called a frequency histogram; one could also make a *density histogram* in which the heights of the bars are scaled appropriately so that the total area of all the bars sums to 1.)

Variation of numerical variables

FIGURE 1.3 depicts a histogram of daily average temperatures in two American cities—San Diego, CA, and Rapid City, SD—for every day from January 1995 to November 2011. Temperature is an example of a *numerical variable*, or something for which numerical comparisons are meaningful (twice as far, six times as fast, \$17 cheaper, and so forth). Numerical variables can be discrete or continuous. Marbles are discrete; we count them on our fingers and toes. Speed and distance are continuous; we measure them in arbitrarily small increments.

Figure 1.3 suggest two obvious, meaningful questions we can ask about a numerical variable like temperature: where is the middle of the sample, and how much does a typical case vary from the middle?

You’re probably already aware of more than more way to answer the question, “Where is the middle?”

- There’s the sample mean, written as \bar{y} . If we have n data

points $\{y_1, \dots, y_n\}$, then

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i.$$

The subscript i 's run from case 1 to case n , where n is the number of data points in the sample. In many data sets the actual ordering of cases won't matter, and will just reflect the arbitrary ordering of the rows in your data frame.⁸

- There's the median, or the halfway point in a sample.
- There's also the mode, or the most common value.

These different ways of quantifying the middle value all have different properties. For example, the median is less sensitive than the mean to extreme values in your sample; there can be more than one mode in a sample, but only one mean or median.⁹

⁸ An obvious exception is in the analysis of time-series data, where the ordering of observations in time may be highly meaningful.

⁹ For example, consider the data set $\{1, 2, 3, 3, 4, 4, 5\}$.

Sample standard deviation and sample variance

Another important question is, "How spread out are the data points from the middle?" Figure 1.3 drives home the importance of dispersion in making useful comparisons. Not only are average temperatures lower overall in Rapid City than in San Diego, but they are also a lot more variable: the coldest days are much colder in Rapid City, but the hottest days are hotter, too.

As with the notion of "middle" itself, there is more than one way of quantifying variability, and each way is appropriate for different purposes. Let's follow the line of thinking that leads us to the *standard deviation*, which is probably the most common way of measuring dispersion. Suppose we choose to measure the middle of a sample y_1, \dots, y_n using the mean, \bar{y} . Each case varies from this middle value by its *deviation*, $y_i - \bar{y}$. Why not, therefore, just compute the average deviation from the mean? Well, because

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y}) &= \frac{1}{n} \sum_{i=1}^n y_i - \frac{n}{n} \bar{y} \\ &= \bar{y} - \bar{y} \\ &= 0. \end{aligned}$$

The positives and negatives cancel each other out. We could certainly fix this by taking the absolute value of each deviation, and then averaging those:

$$M = \frac{1}{n} \sum_{i=1}^n |y_i - \bar{y}|.$$

This quantity is a perfectly sensible measure of the “typical deviation” from the middle. Fittingly enough, it is called the *mean absolute deviation* of the sample.

But it turns out that, for the purposes of statistical modeling, a quantity called the *sample variance* makes more sense:

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2.$$

That is, we *square* each deviation from \bar{y} , rather than take the absolute value. Remember that when we square a negative number, it becomes positive, so that we don’t have the problem of the positives and negatives cancelling each other out.

Two natural questions about this quantity are:

- (1) Why do we divide by $n - 1$, when dividing by n would seem to make more sense for computing an average?
- (2) Why do we square the deviations, instead of taking absolute values as above?

To answer the first question: we divide by $n - 1$ rather than n for obscure technical reasons that, despite what you may read in other statistics textbooks, just aren’t that important. (It has to do with “unbiased estimators,” which, despite the appealing name, are overrated.) Mainly we use $n - 1$ to follow convention.

As for the second question: because sums of squares are special! In all seriousness, there are deep mathematical reasons why we choose to measure dispersion using sums of squared deviations, rather than the seemingly more natural sums of absolute deviations. You’ll learn why in the next chapter, but if you want a preview, think about Pythagoras and right triangles. . . .

Of course, computing the sample variance leaves us in the awkward position of measuring variation in the *squared* units of whatever our variable is measured in. This is not intuitive; imagining telling someone that the mean temperature in Rapid City over the last 17 years was 47.3 degrees Fahrenheit, with a sample variance of 402 degrees squared. This is a true statement, but nearly uninterpretable.

Luckily, this is easily fixed by taking the square root of the sample variance, giving us the sample standard deviation:

$$s = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2}.$$




Figure 1.4: The histogram shows average daily temperatures in Rapid City. The blue dot is the sample mean, and the blue line shows an interval encompassing one sample standard deviation to either side of the sample mean.

Now we’re back to the original units, and an interpretable measure of “typical deviation from the middle”—for Rapid City, 20.1 degrees. This looks about right from the histogram above; the blue dot is the sample mean, and the blue line stretches 1 sample standard deviation to either side of the mean.

Two other simple measures of spread are worth mentioning briefly. First, there’s the *range*, or the difference between the largest and smallest values in the sample. There’s also the *interquartile range*, or the difference between the 75th and 25th percentiles. This is robust to extreme values, since it involves only the middle 50% of the sample.

Percentiles, quantiles, and coverage intervals

Another useful way to summarize the variation of a numerical variable across cases is to compute a set of percentiles, also called quantiles. A familiar example is the median: it happens that exactly 50% of the daily average temperatures in Rapid City fall below 47.6 degrees, and we call this point the median (or the 50th percentile). Similarly, 10% of days in Rapid City are colder than 20.7 degrees, and 90% of days are colder than 73.2 degrees; these are the 10th and 90th percentiles, respectively. A quantile is just a percentile expressed in terms of a decimal fraction; the 80th percentile and 0.8 quantile are the same number.

A common way to summarize a distribution of a numerical variable is to quote a *coverage interval* defined by two percentiles, like the 10th and 90th percentiles (which covers 80% of the cases) or the 2.5th and 97.5th percentiles (which covers 95% of the cases). So, for example, we might quote an 80% coverage interval for daily average temperatures in Rapid City as (20.7, 73.2), whose endpoints are formed from the 10th and 90th percentiles.

Variation between, and within, groups

A COMMON situation is that we have both categorical and numerical data about each case in a data set. For example, Table 1.3 below shows the average SAT math and verbal scores, stratified by college, for undergraduates in the incoming fall of 2000 freshmen class at the University of Texas at Austin. All 5,191 students who went on to receive a bachelor's degree within 6 years are included; those who dropped out, for whatever reason, are not.

College	Average SAT	
	Math	Verbal
Architecture	685	662
Business	633	597
Communications	592	609
Education	555	546
Engineering	675	606
Fine Arts	597	594
Liberal Arts	598	590
Natural Sciences	633	597
Nursing	561	555
Social Work	602	589

Table 1.3: Average SAT math and verbal scores, stratified by college, for entering freshmen at UT–Austin in the fall of 2000. Collected under the Freedom of Information Act from the state of Texas.

The table tells you something about how the numerical variables (test scores) change depending upon the categorical variable (college), and they are superficially similar to the contingency tables we just encountered. They highlight interesting and useful facts about variation between the groups. Math skills, for example, are probably more important for engineering majors than English majors, and this is reflected in the differences between the group-level means.

Table 1.3 does differ from a contingency table, however, in one crucial respect: the entries in the table are not counts, but group-level averages. Notice that, to depict between-group variation, the table has reduced each college to a typical case, represented by some hypothetical student who earned the college-wide average SAT scores on both the math and verbal sections. In doing so, it has obscured the underlying variability of students *within* the colleges. But as our example of city temperatures demonstrated, sometimes this variability is an important part of the story as well.

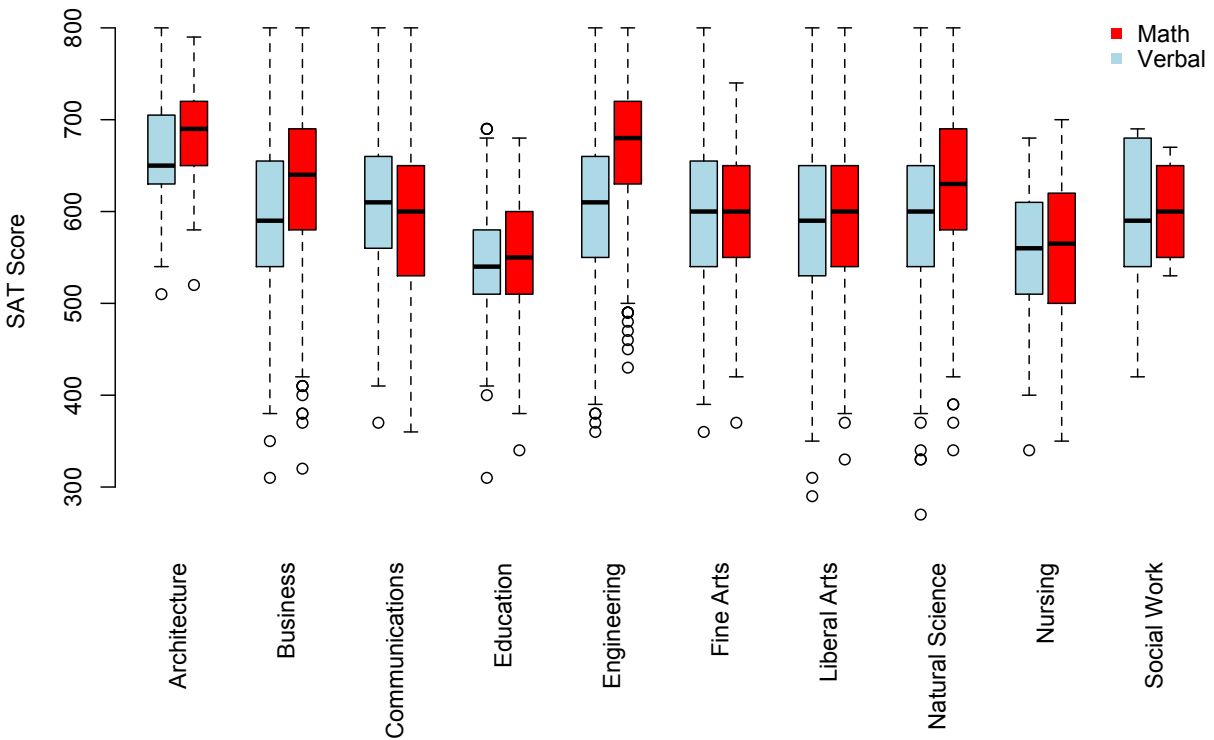


Figure 1.5: Boxplots of the full data set used to form the means in Table 1.3. Each box shows the interquartile range for a specified subset of the data; the middle line within each box is the median of that subset. In this boxplot, the whiskers extend outside the box no further than 1.5 times the interquartile range. Points outside this interval are shown as individual dots.

Boxplots

This is where boxplots are useful: they allow you to assess variability both between and within the groups. Many times it is the within-group variability that matters most. For example, as Figure 1.5 shows, SAT scores vary at least as much within a college as they do between colleges. For example, the 52-point difference in average SAT math scores between Architecture students and Natural Science students looks a little smaller compared to the roughly 100-point interquartile range of math scores within Natural Science. Although between-group variability cannot be ignored, it is nonetheless dwarfed by within-group variability.

The situation is quite different Figure 1.6. These boxplots show the growth of guinea pigs’ teeth versus their daily dosage of Vitamin C. Like humans, but unlike most other mammals, guinea pigs need Vitamin C to survive, yet cannot synthesize their own. The amount they are given predicts much about their health, measured in this case by the length of their teeth. We therefore see compara-




Figure 1.6: For comparison, the table of within-group means is below. Notice how the within-group variability evident in the boxplots at left simply disappears when presented in the form of a summary table, below.


Dose (mg)	Tooth len.
0.5	7.98
1.0	16.77
2.0	26.14

tively more variability between the groups, whose boxplots almost don't overlap.

The same comparison will come up again and again: between-group variability (the differences between typical or average group members) versus within-group variability (the variation of cases within a single group). We'll soon make this comparison mathematically rigorous, but these examples convey the essence of the idea:

- A UT student's college tells you something, though not everything, about his or her likely SAT scores.
- A guinea pig's Vitamin C regimen tells you something, though not everything, about its tooth growth. But in a relative sense, it tells you more than a UT student's college tells you about his or her SAT scores.

Always remember that a table of group-wise means does not depict "data" as such, but an abstraction of some typical group member. This abstraction may be useful for some purposes. But within-group variability is also important, and may even be the dominant feature of interest. In this case, presenting the group-wise means alone, without the corresponding plots or measures of variability, may obscure more than it reveals.




Dot plots

The *dot plot*, or strip chart, is a close cousin of the boxplot. For example, the plot above depicts a relationship between the length of a mammal's dreams and the severity of the danger it faces from predators. Each dot is a single species of mammal—like, for example, the critter at right. The predation index (X) is an ordinal variable running from 1 (least danger) to 5 (most danger). It accounts both for how likely an animal is to be preyed upon, and how exposed it is when sleeping. Notice the direction of the trend—you'd sleep poorly too if you were worried about being eaten.

As you can see, the dot plot is useful for small data sets, when a boxplot is no simpler than just plotting the cases group by group. Strictly speaking, the points should all line up vertically with their corresponding values of predation index, on the x -axis. But a bit of meaningless horizontal jitter has been added to the dots, which allows the eye to distinguish the individual cases more easily.

Figure 1.7: Dreaming hours per night versus danger of predation for 50 mammalian species. Data from: "Sleep in Mammals: Ecological and Constitutional Correlates," Allison and Cicchetti (1976). *Science*, November 12, vol. 194, pp. 732-734. Photo of the dreaming critter from the MIT News office (web.mit.edu/newsoffice/2001/dreaming.html).






Figure 1.8: Daily peak electricity demand (stratified by month) in Raleigh, NC from 2006–09. The dashed line is the average peak demand for the whole data set, and the blue dots are the month-by-month means.

Dot plots can also be effective for larger data sets. Here we see four years of data on daily peak electricity demand for the city of Raleigh, NC, stratified by month of the year. Both the between-group and within-group variation show up clearly.

Group means and grand means

If you looked carefully, you may have noticed two extra features of the dot plots in Figures 1.7 and 1.8. The square blue dots show the *group means* for each category. The dotted green line shows the *grand mean* for the entire data set, irrespective of group identity. Notice that, in plotting these means along with the data, we have implicitly partitioned the variability:

$$\begin{aligned}
 \text{Individual case} &= \text{Group mean} + \text{Deviation of that case} & (1.1) \\
 &= \text{Grand mean} + \text{Deviation of group} + \text{Deviation of case} .
 \end{aligned}$$

This is just about the simplest statistical model we can fit, but it's still very powerful. We'll revisit it soon.




Figure 1.9: A scatter plot of the daily returns for Microsoft stock, versus those of Apple stock, for every trading day in 2015. The daily return is the implied interest rate from holding a stock from the end of one trading day to the end of the next. For example, Apple stock closed at \$105.95 per share on January 7th and at \$110.02 on January 8th. Thus the return for January 8th was

$$\frac{110.02 - 105.95}{105.95} \approx 0.038,$$

or about a 3.8% daily return. On the same day, holders of Microsoft stock enjoyed a 2.9% return.

More than one numerical variable

Our basic tool for visualizing a bivariate relationship between two numerical variables is the *scatter plot*. Figure 1.9 shows a plot of the daily returns for Microsoft stock versus Apple stock for every trading day in 2015. Every dot corresponds to a day. The location of the dot along the horizontal axis shows the Apple return, and the location on the vertical axis shows the Microsoft return, for that day. In this case, we can see that Microsoft and Apple stocks tend to move up and down together. (Most stocks do.) We can also see the speckling of outliers: those points that are visibly separate from the main cloud and that represent very good (or bad) days for holders of these two stocks.

A simple way to visualize three or more numerical variables is via a *pairs plot*, as in Figure 1.10. A pairs plot is a matrix of simpler plots, each depicting a bivariate relationship. In Figure 1.10, we see scatterplots for each pair of the daily returns for Microsoft, Facebook, Apple, and Amazon stocks. The histograms on the diagonal serve a dual purpose: (1) they show the variability of each stock in isolation; and (2) they label the rows and columns, so that you know which plots compare which variables.

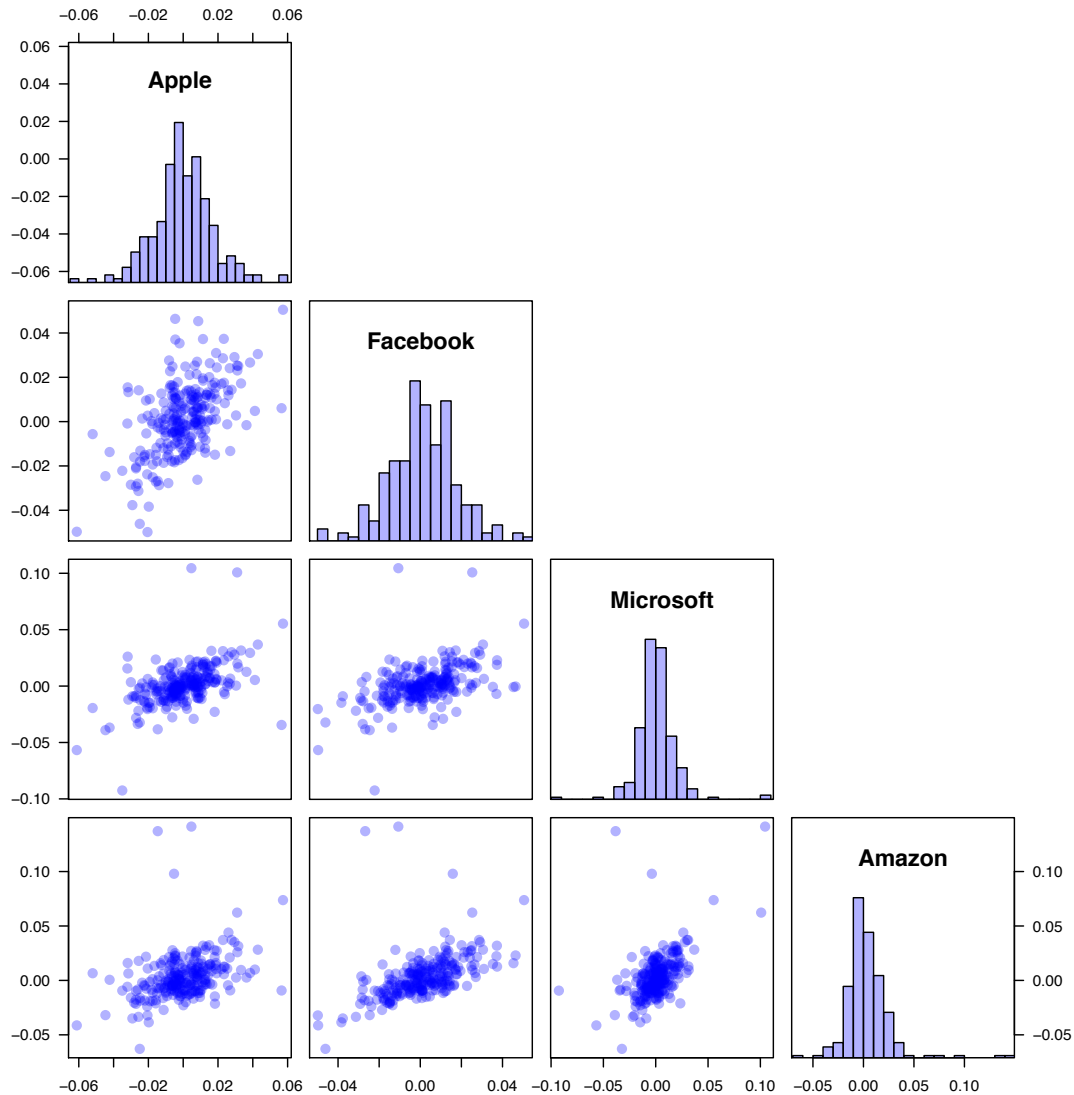


Figure 1.10: A pairs plot: a matrix of four pairwise scatter plots for the daily returns of Apple, Facebook, Microsoft, and Amazon stocks in 2015. The histograms along the diagonal also label the rows and columns of the matrix: e.g. the plot in the second row has Facebook returns along the vertical axis, while the plots in the second column both have Facebook returns along the horizontal axis.

Sample correlation. The *sample correlation coefficient* is a standard measure of the strength of linear dependence between two variables in a sample. If we label the first variable as x_1, \dots, x_n and the second as y_1, \dots, y_n , then the correlation coefficient is defined as

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{(n-1)s_x s_y}, \quad (1.2)$$

where s_x and s_y are the sample standard deviations of the X and Y variables. At right you see scatter plots that depict examples of strong positive (top) and weak negative (bottom) correlation. Sample correlation is between 1 and -1 , which are the extremes of perfect positive and perfect negative correlation.


To summarize the correlation among a set of more than two variables, we typically calculate a *correlation matrix* whose entry in row i , column j is the correlation between variable i and variable j . For the four stocks depicted in Figure 1.10, the correlation matrix R is below. Notice that the matrix is symmetric and has ones along the diagonal (because a variable is perfectly correlated with itself):

	Apple	Microsoft	Facebook	Amazon
Apple	1.00	0.52	0.55	0.36
Microsoft	0.52	1.00	0.47	0.52
Facebook	0.55	0.47	1.00	0.50
Amazon	0.36	0.52	0.50	1.00

Caveats. A key fact to remember is that correlation measures the strength of *linear* dependence; if two variables don't fall roughly along a straight line in a scatter plot, then correlation can be misleading. Therefore, always plot your data, as in Figure 1.11: four different data sets, four different stories about what's going on. Yet all have the same correlation: $r = 0.816$.

Another important fact is that a sample correlation of 0 ("uncorrelated") does not necessarily mean that two variables are unrelated. In fact, the correlation coefficient is so intimately tied up with the assumption of a linear relationship that it breaks down entirely when used to quantify the strength of nonlinear relationships. In each of the six plots in Figure 1.12, for example, there is an obvious nonlinear relationship between the two variables. Yet the sample correlation coefficient for each of them is exactly zero.

The moral of the story is that sample correlation is just one number, and can only tell you so much about the relationship between two sets of numbers.







Figure 1.11: above. Data taken from F.J. Anscombe, "Graphs in Statistical Analysis." *American Statistician*, 27 (1973), pp. 17–21

Figure 1.12: left. Each panel shows obvious dependence, but has a sample correlation of $r = 0$.






Figure 1.13: Highway gas mileage versus engine power for 387 vehicles in five different classes.

Further multivariate plots

In visualizing data, we are constrained by the limitations of the two-dimensional page or screen. Nevertheless, there are many techniques for showing more than two variables at once, despite these limitations.

Lattice plots

Figure 1.13 shows three variables from a data set on 387 vehicles: the highway gas mileage, the engine power (in horsepower), and the class of the vehicle (minivan, sedan, sports car, SUV, or wagon). This is done via a *lattice plot*, which displays a variable or the relationship between two variables, stratified by the value of some other variable. In this case the main relationship of interest is between mileage and engine power, and the stratifying variable is vehicle class. Notice how figure 1.13 repeats a scatterplot of MPG versus horsepower five times: one plot for the vehicles in each class. To facilitate comparisons across the strata, both the horizontal and vertical axes are identical in each plot.

The figure suggests several facts:

- For a fixed level of engine power, there is considerable variability in fuel economy. (Pick a fixed point on the horizontal axis and focus on the cars near there. Now look at the corresponding variability along the vertical axis for those cars.)
- The overall MPG–horsepower trend is negative for all classes.
- The SUVs have the worst gas mileage overall, and in particular have worse mileage than the sports cars and wagons

Another term for a lattice plot is a trellis plot.




Figure 1.14: Highway gas mileage versus engine power for 387 vehicles in five different classes.

despite having similar or lower power. (Compare the average vertical location in the SUV panel versus the others).

- The MPG–horsepower relationship becomes nonlinear for sedans at low horsepower, but perhaps not for wagons.
- As engine power increases, the dropoff in gas mileage looks steeper for SUVs than for sports cars.
- Nobody makes a powerful minivan.

We can make a lattice of boxplots as well. For example, Figure 1.14 shows boxplots of engine power versus number of **engine cylinders**, stratified by vehicle class. This suggests an explanation for the fact that engine power is not a perfect predictor of fuel economy: some cars get more power out of a smaller engine, and are presumably more efficient as a result.

With a numerical variable. In Figure 1.13, the stratifying variable is categorical. But we can also stratify a data set according to a numerical variable, by discretizing that variable into bins—much in the same way we do when we make a histogram. Figure 1.15 shows the latitude, longitude, and depth (in kilometers) beneath the earth’s surface for the epicenter of every earthquake recorded since 1963 near Fiji, an island in the South Pacific Ocean. The depth variable has been stratified into nine equal-length intervals. The nine panels show the latitude and longitude of the quakes whose depths fell in each interval, labeled at the top of each panel.

As depth increases (going left to right, top to bottom), a spatial pattern emerges. The shallower earthquakes are at the intersection of two major tectonic plates. The deeper quakes emanate from the Tonga Trench—35,702 feet below the sea at its deepest point.¹⁰

¹⁰ And the final resting place of 3.9 kilograms of radioactive plutonium-238 from the ill-fated Apollo 13 mission.

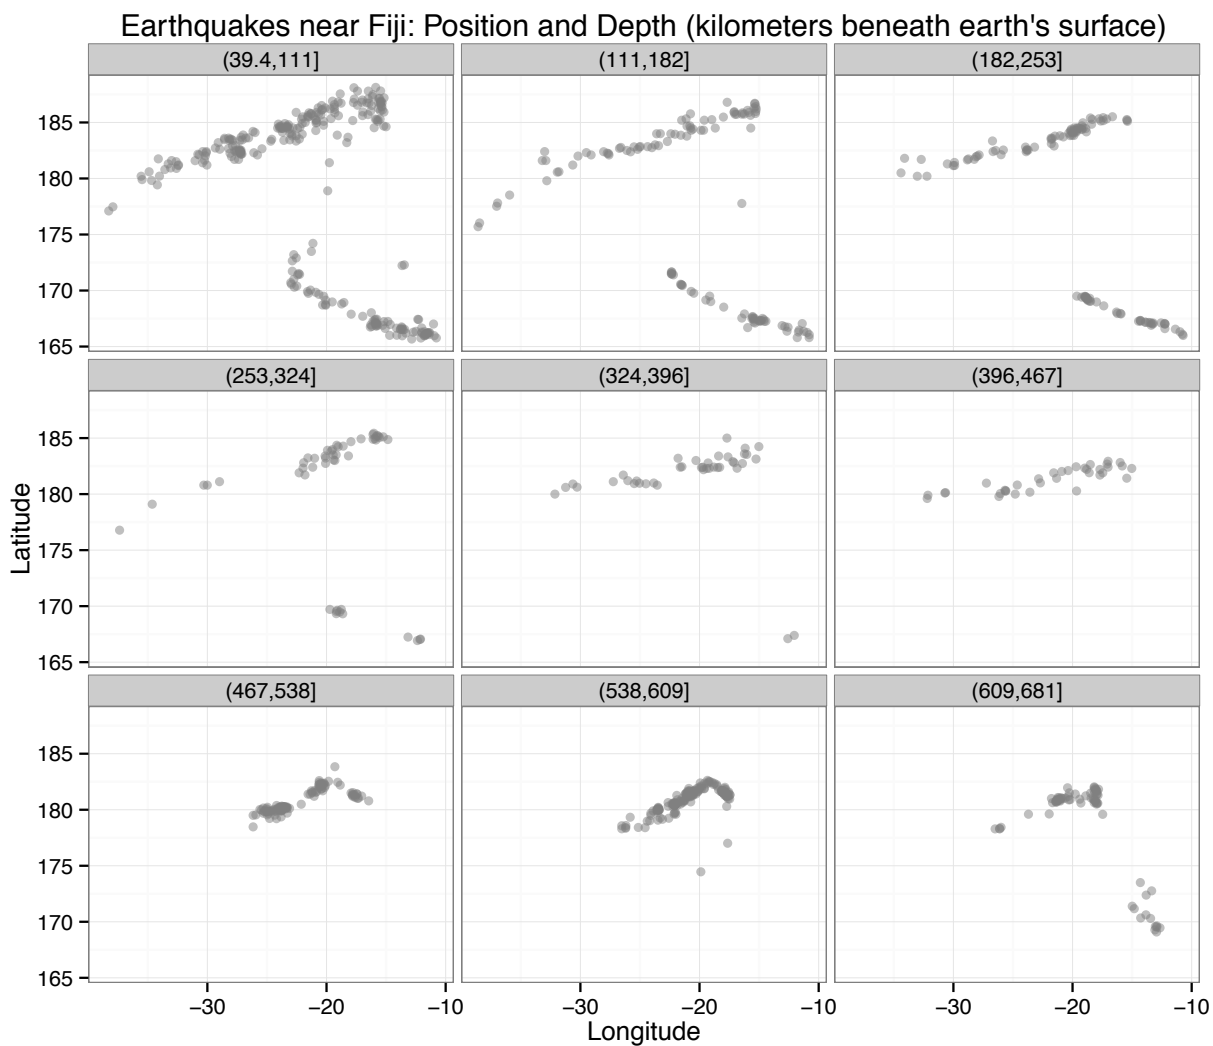


Figure 1.15: Earthquakes in Fiji: latitude versus longitude for quakes within each of nine different depth intervals. Here the range of depths beneath the earth's surface (in kilometers) is labeled at the top of each panel.